



# LINEAR REGRESSION

Dr. Ram Prasad K  
VisionCog R&D

[ram.krish@visioncog.com](mailto:ram.krish@visioncog.com)  
<https://www.visioncog.com>

# MACHINE LEARNING



## Machine Learning:

An algorithmic way of *making sense (learning) from data*.

## Applications:

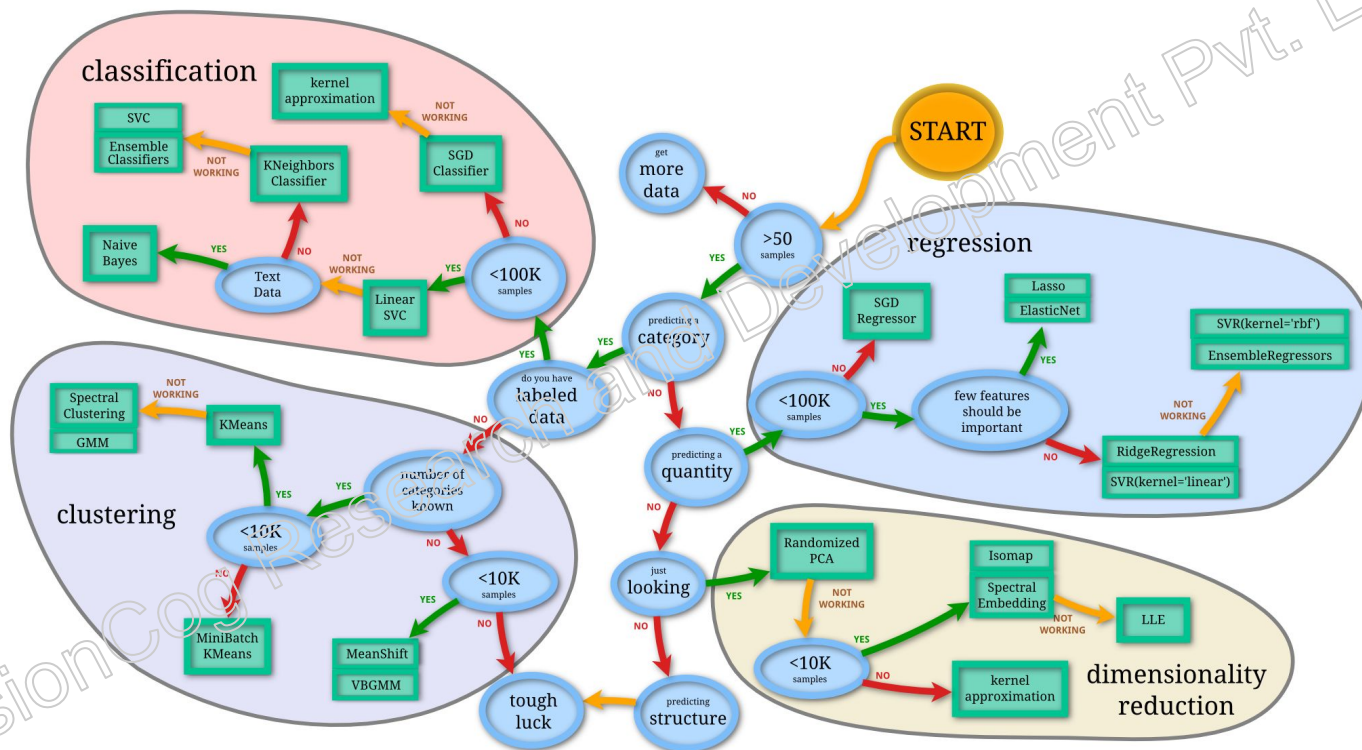
- Spam filters (**Classification**)
- Predict height based on weight and age (**Regression**)
- Online recommendation systems (**Clustering**)
- Visualizing multidimensional data (**Dimensionality reduction**)

# MACHINE LEARNING

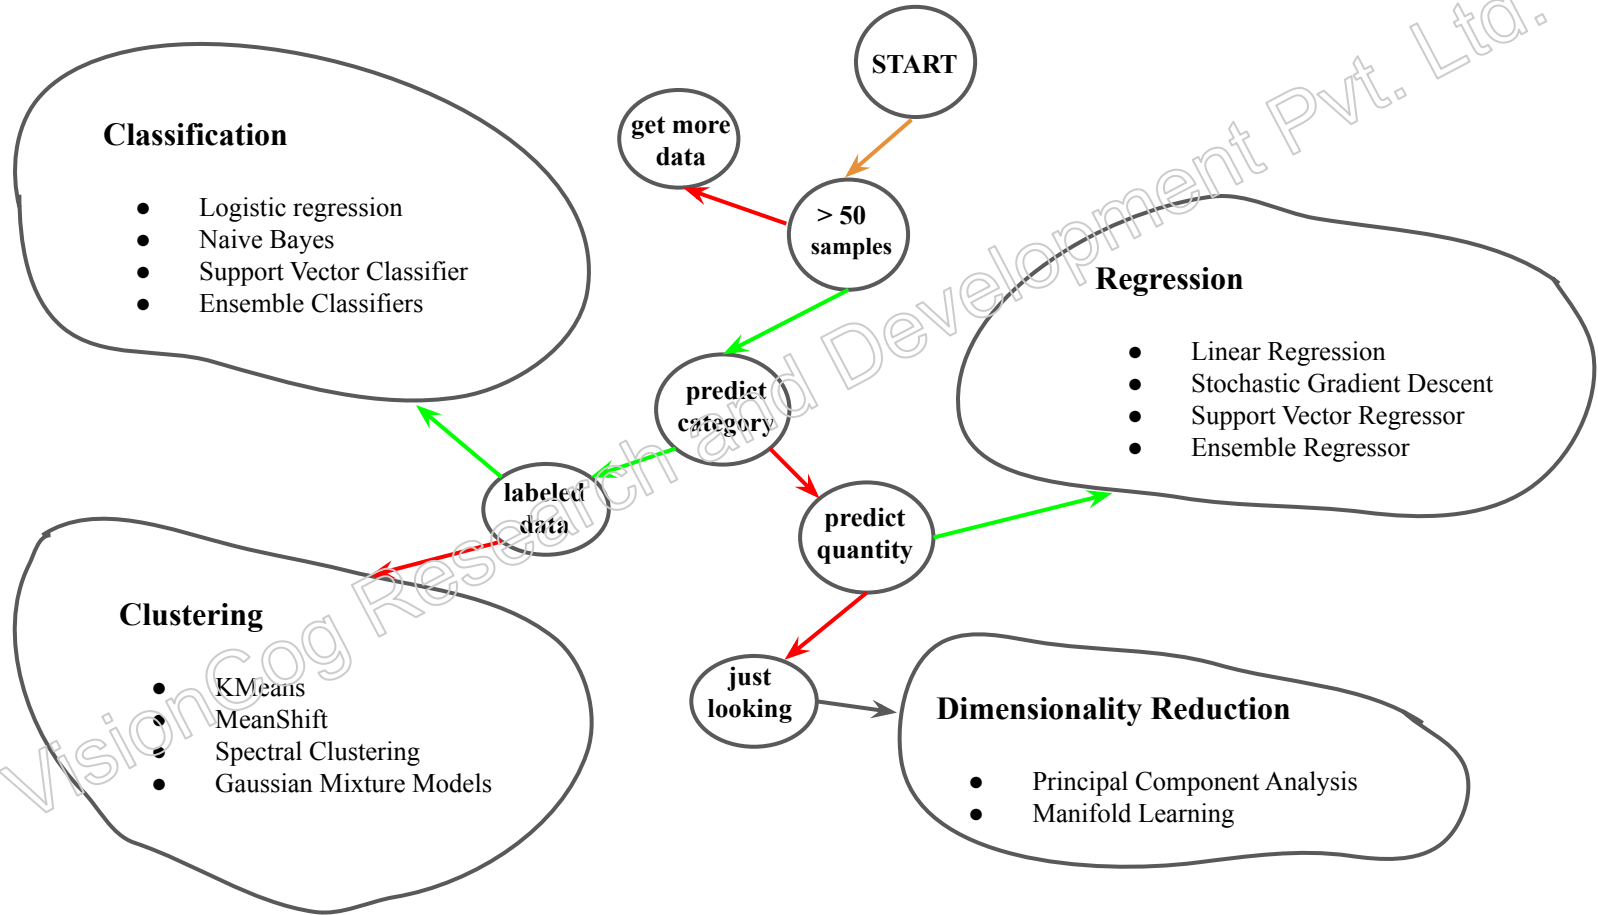


## Scikit Learn

- Machine Learning library - in **Python**
- Simple and efficient tools for data analysis
- Built on NumPy, SciPy, and matplotlib
- API is remarkably well designed



# MACHINE LEARNING





# LINEAR REGRESSION

# LINEAR REGRESSION



## Dependent and Independent variable

Expression	Independent	Dependent
$y = 3 + 2x$	$x$	$y$
$y = x^2 - 2x$	$x$	$y$
$z = 5x^2 + 8y^3$	$x, y$	$z$

## Regression:

Modeling a relationship between *dependent* and *independent* variables for *prediction*.

# LINEAR REGRESSION



***Simple Linear Regression*** or ***Univariate Linear Regression***.

Only one independent variable

***Multiple Linear Regression*** or ***Multivariate Linear Regression***.

More than one independent variable





# LINEAR REGRESSION



```
import numpy as np
np.random.seed(42)
```

```
X = 2 * np.random.rand(100, 1)
```

```
y = 4 + 3 * X + np.random.randn(100, 1)
```



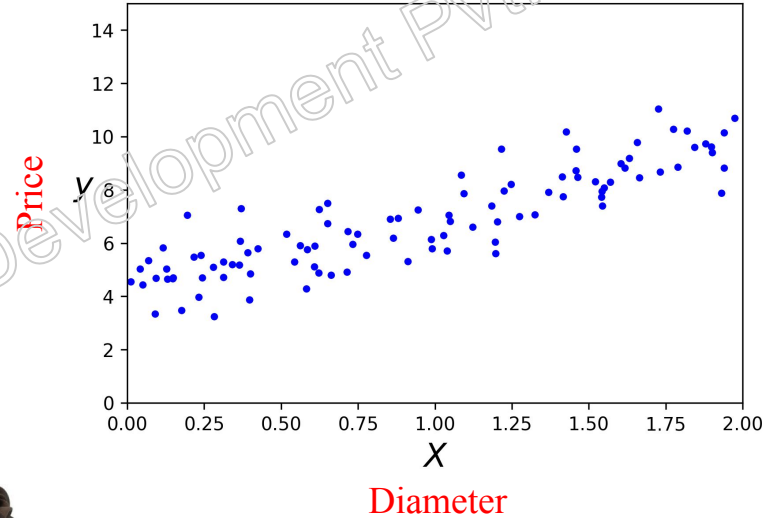
# LINEAR REGRESSION



```
import numpy as np
np.random.seed(42)

X = 2 * np.random.rand(100, 1)

y = 4 + 3 * X + np.random.randn(100, 1)
```



# LINEAR REGRESSION




```
import numpy as np  
np.random.seed(42)
```

```
X = 2 * np.random.rand(100, 1)
```

```
y = 4 + 3 * X + np.random.randn(100, 1)
```

Now assume you don't know how **y** was calculated from **X**



# LINEAR REGRESSION



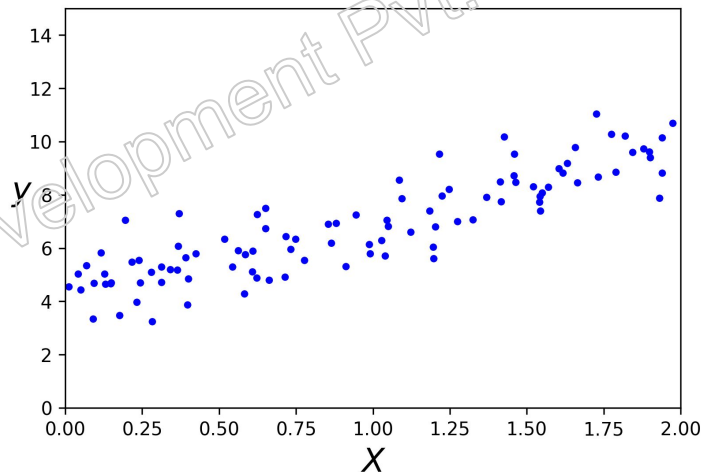
Assume you now only have **X** and **y**.

Split the dataset into training and test set.

Now assume you don't know how **y** was calculated from **X**

```
from sklearn.model_selection import train_test_split
```


```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size = 0.20, random_state = 42)
```




# LINEAR REGRESSION




**80% training**





**20% testing**



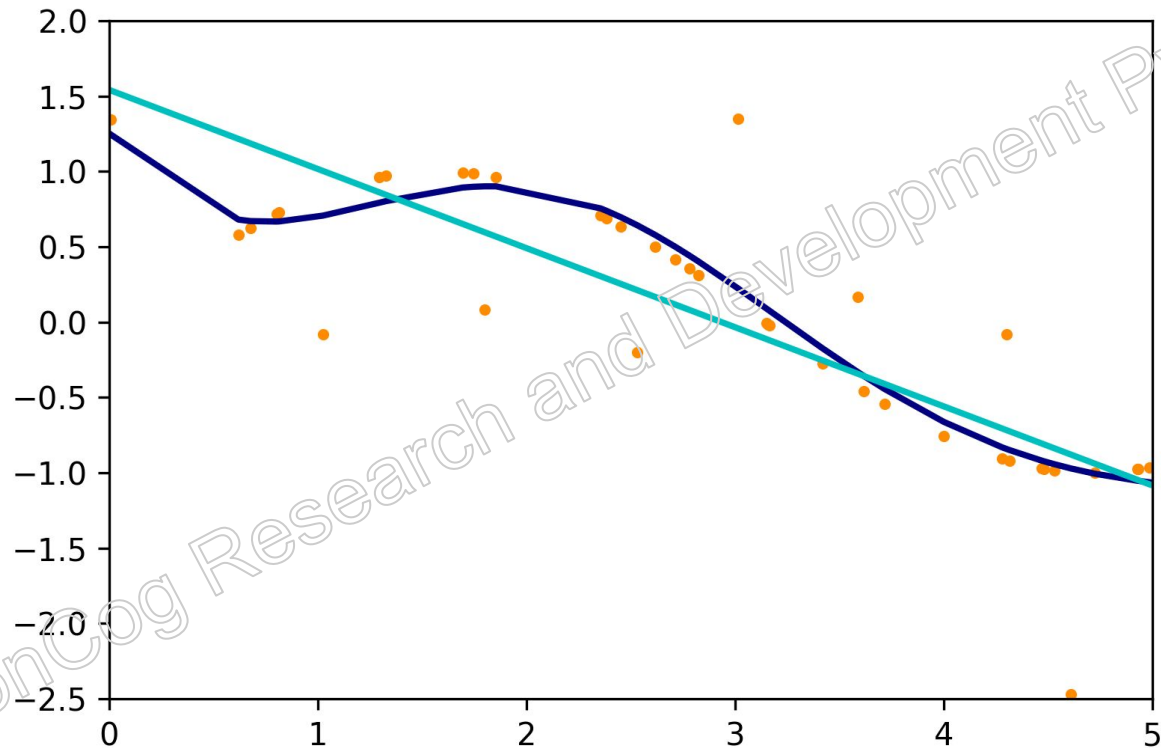
# LINEAR REGRESSION



# LINEAR REGRESSION





# LINEAR REGRESSION







# LINEAR REGRESSION



# LINEAR REGRESSION



# LINEAR REGRESSION



We can obtain here a **straight line** which passes close to as many points as possible.

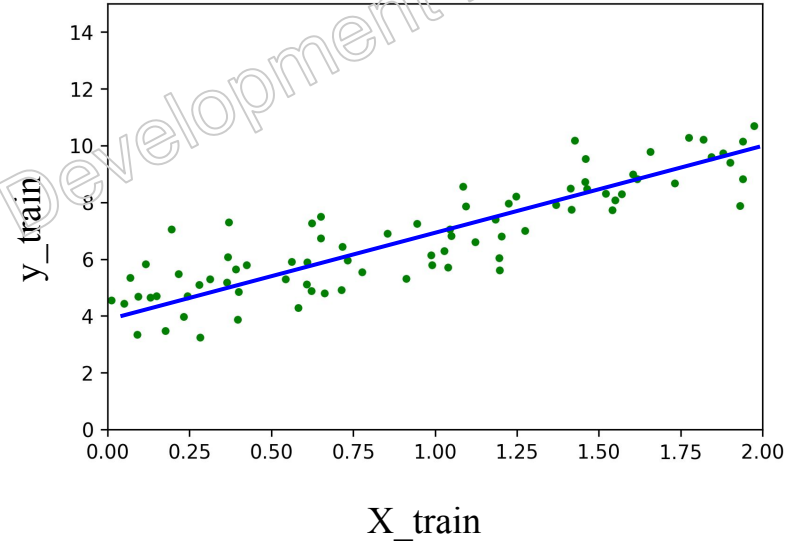
What parameters are required to represent a straight line?

y-intercept

slope

Equation of a straight line:

$$y = c + mx$$



# LINEAR REGRESSION




Mathematical model for Simple Linear Regress:

$$\hat{y} = \theta_0 + \theta_1 x$$

Diagram illustrating the components of the linear regression equation:

- $\theta_0$  is labeled as the **intercept** (in red text).
- $\theta_1$  is labeled as the **coefficient** (in green text).
- Both  $\theta_0$  and  $\theta_1$  are collectively labeled as **model parameters** (in blue text).



*The line models the relationship between cake independent and dependent variable.*

# LINEAR REGRESSION



## General/Multiple Linear Regression

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$\hat{y}$  is the predicted value

$n$  is the number of features

$x_i$  is the  $i^{th}$  feature value

$\theta_j$  is the  $j^{th}$  model parameter

$\theta_0$  is the intercept (also called **bias** term)

# LINEAR REGRESSION



## Vectorized general form

$$\hat{y} = h_{\theta}(\mathbf{x}) = \theta^T \cdot \mathbf{x}$$

$\theta$  is the models **parameter** vector

$\theta_0$  is the *bias/intercept*

$\theta_1, \theta_2, \dots, \theta_n$  are **coefficients** or feature weights.

$\mathbf{x}$  is the **feature** vector  $x_0$  to  $x_n$  with  $x_0$  always 1

$\theta^T \cdot \mathbf{x}$  is the dot product of  $\theta^T$  and  $\mathbf{x}$

$h_{\theta}$  is the **hypothesis** function using model parameters  $\theta$

# LINEAR REGRESSION



$$\hat{y} = h_{\theta}(\mathbf{x}) = \theta^T \cdot \mathbf{x}$$

$$\text{MSE}(\mathbf{X}, h_{\theta}) = \frac{1}{m} \sum_{i=1}^m (\theta^T \mathbf{x}^{(i)} - y^{(i)})^2$$

cost function

## Normal Equation


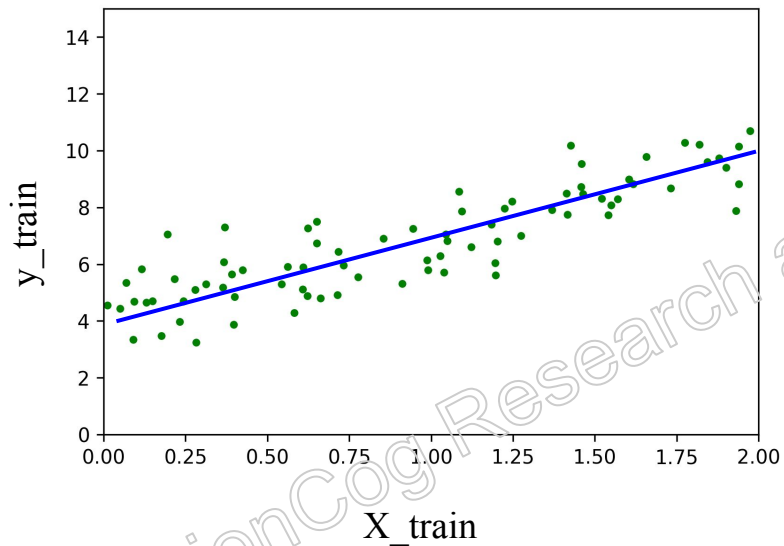
To find the parameters, we have a closed-form solution:

$$\hat{\theta} = (\mathbf{X}^T \cdot \mathbf{X})^{-1} \cdot \mathbf{X}^T \cdot \mathbf{y}$$

$\hat{\theta}$  is the value of  $\theta$  that minimizes the cost function (least squares)

$\mathbf{y}$  is the vector of target values

# LINEAR REGRESSION





# LINEAR REGRESSION




Mathematical model for Simple Linear Regress:

$$\hat{y} = \theta_0 + \theta_1 x$$

Diagram illustrating the components of the linear regression equation:

- $\theta_0$  is labeled as the **intercept** (in red text).
- $\theta_1$  is labeled as the **coefficient** (in green text).
- Both  $\theta_0$  and  $\theta_1$  are collectively labeled as **model parameters** (in blue text).



*The line models the relationship between cake independent and dependent variable.*


# LINEAR REGRESSION



```
import numpy as np
np.random.seed(42)
```

```
X = 2 * np.random.rand(100, 1)
```

```
y = 4 + 3 * X + np.random.randn(100, 1)
```



```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.20, random_state = 42)
```

# LINEAR REGRESSION



```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
print(model.intercept_)
```

```
print(model.coef_)
```

```
# [4.14291332]
```

```
# [[2.79932366]]
```

$$y = 4 + 3x + \varepsilon$$

$$y = 4.14 + 2.78x$$

VisionCog Research and Development Pvt. Ltd.

# LINEAR REGRESSION



## Evaluating the model

### Testing error

**$R^2$  / Coefficient of determination / Goodness-of-fit**

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

$$SS_{res} = \sum_i^n (y_i - f(x_i))^2$$

sum of squares explained by model

$$SS_{tot} = \sum_{i=1}^n (y_i - \bar{y})^2$$

sum of squares around the mean

# LINEAR REGRESSION



```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
print(model.intercept_)
```

```
print(model.coef_)
```

```
# [4.14291332]
```

```
# [[2.79932366]]
```

```
score = model.score(X_test, y_test)
```

```
print(score)
```

```
# 0.8072059636181392
```

$$y = 4 + 3x + \varepsilon$$

$$y = 4.14 + 2.78x$$



# MODEL EVALUATION USING CROSS-VALIDATION

# LINEAR REGRESSION








## Cross-validation

- Training dataset is further divided into training and validation.
- The validation splitting is based on the number of folds.
- k-Fold cross validation divides the training set into k equal blocks.
- In each round:
  - one of the blocks is used as validation set, and
  - the remaining is used as training set.

Training set



Divided into k blocks (k = 10)



# LINEAR REGRESSION



```
from sklearn.model_selection import cross_val_score
```

```
scores_CV = cross_val_score(model, X_train, y_train, cv=10)
```

```
print(scores_CV)
```

```
# [0.88261122 0.83015975 0.5121498 0.57767211 0.21819067 0.80309185
```

```
# 0.73759195 0.38341108 0.77308338 0.784341 ]
```

```
print(scores_CV.mean())
```

```
# 0.6502302819222789
```

```
print(scores_CV.std())
```

```
# 0.2079986593933312
```

**Cross-validated model score: 0.65 +/- 0.21**