



# CONVOLUTIONAL NEURAL NETWORKS (CNN)

Dr. Ram Prasad K  
VisionCog R&D

[ram.krish@visioncog.com](mailto:ram.krish@visioncog.com)  
<https://www.visioncog.com>

# CONVNETS



- DNN can learn only global features.
- They are not robust enough to capture large variations.
- To overcome this issue, we can increase the training set with several variations.
- DNN works reasonably well for small images only.

# CONVNETS

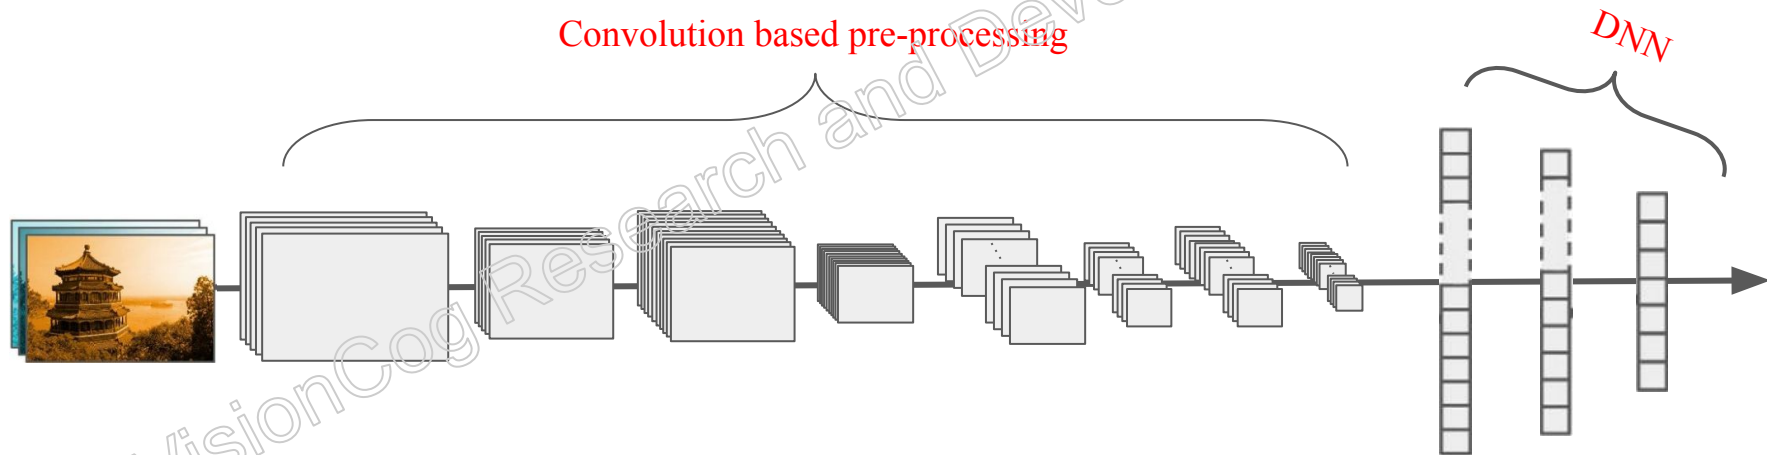


- In computer vision applications, we usually have to deal with complex variations and large images.
- Successful extraction of salient features is key for any computer vision task.
- A model which can learn local features will be robust enough in extracting salient features.
- Convolutional Neural Networks (CNN or ConvNets)
  - DNN modified to capture local features

# CONVNETS



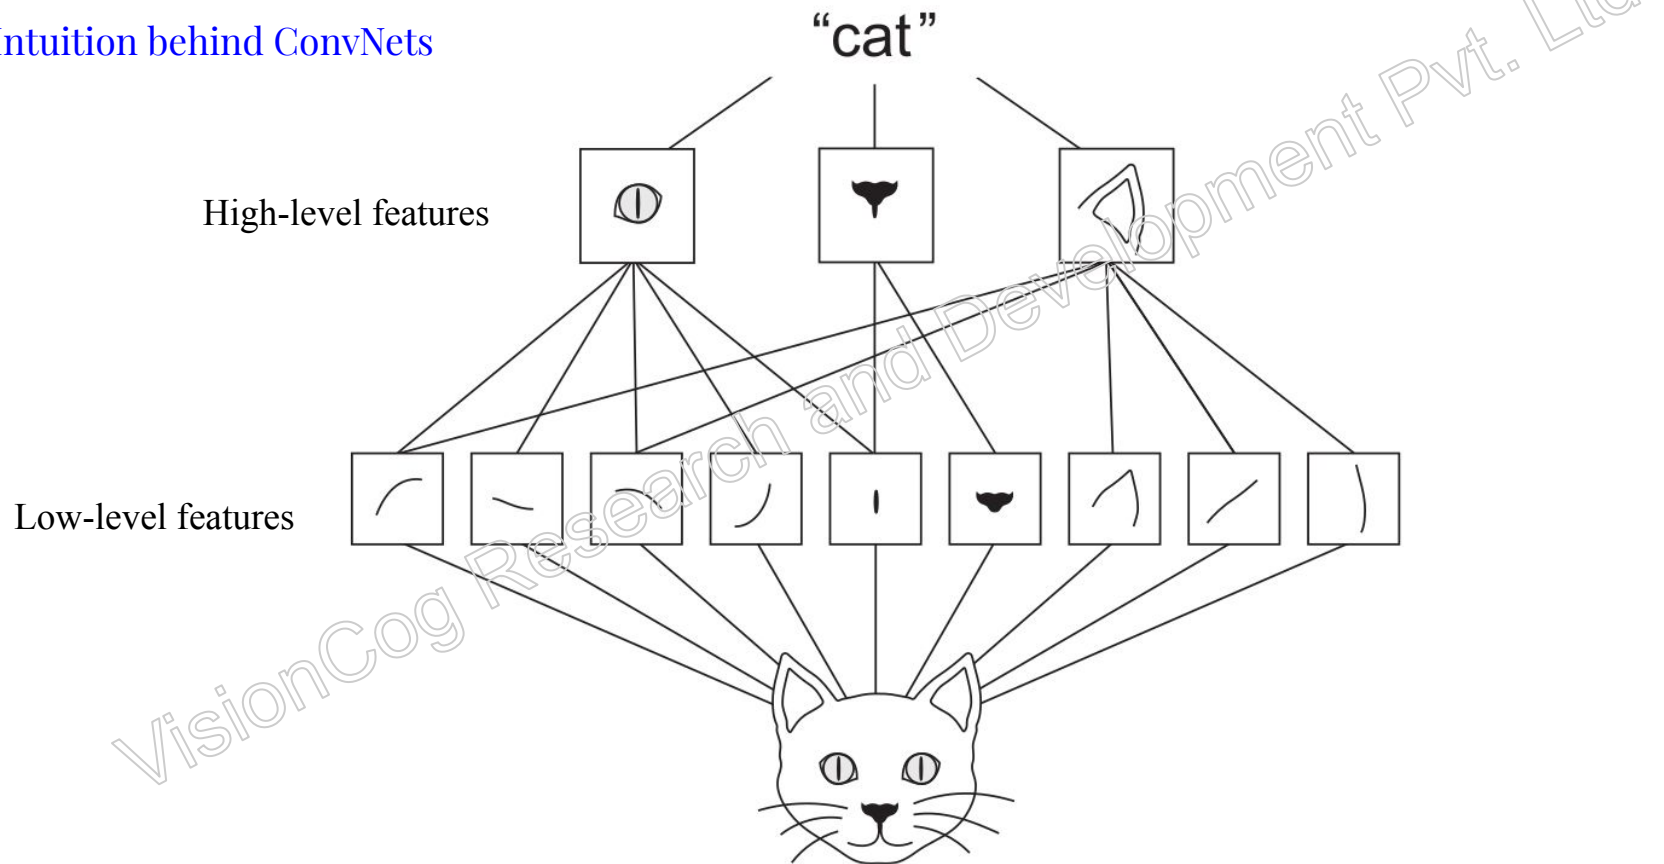
## Intuition behind ConvNets



# CONVNETS



## Intuition behind ConvNets



# CONVNETS



## Building blocks of CNN

### **Convolutions**

Mathematical operation which slides one function over the other and measures integral of pointwise multiplication (i.e., weighted sum of the inputs).

### **Strides**

How quickly window slides.

Stride 2 means, window moves by 2 pixels at a time.

### **Pooling**

Downsampling feature maps.

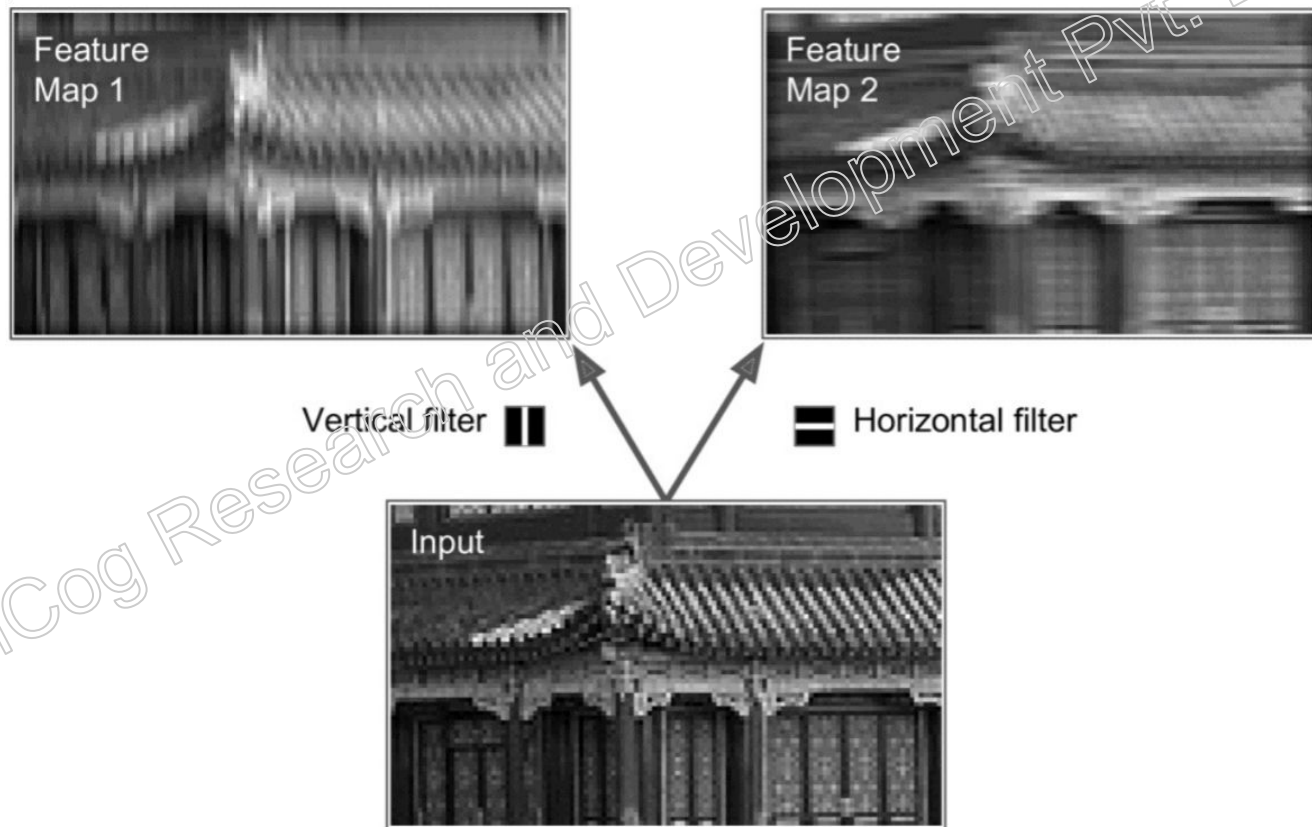
### **DNN**

Fully connected DNN for classification

# CONVNETS



## Convolutions



# CONVNETS



1	0	0	0	0	0
0	1	1	1	0	0
0	0	1	0	0	1
0	1	0	0	1	0
0	0	0	1	0	1
0	0	1	0	0	1

convolution  
stride=1

1	1	2	0
1	3	0	0
2	0	0	3
0	0	3	0

0	0	1
0	1	0
1	0	0

: kernel

0	0	0	1	0	0
0	1	0	0	0	0
0	0	1	1	1	0
0	0	0	1	0	0
0	0	1	0	0	1
0	0	0	0	1	0

convolution  
stride=1


1	1	1	1
0	1	1	2
1	1	3	0
0	2	0	0



# CONVNETS



## Pooling - MaxPooling



# CONVNETS



1	0	0	0	0	0
0	1	1	1	0	0
0	0	1	0	0	1
0	1	0	0	1	0
0	0	0	1	0	1
0	0	1	0	0	1

convolution  
stride=1

1	1	2	0
1	3	0	0
2	0	0	3
0	0	3	0

pooling  
stride=2

3	2
2	3

0	0	1
0	1	0
1	0	0

: kernel

max

.	.
.	.

0	0	0	1	0	0
0	1	0	0	0	0
0	0	1	1	1	0
0	0	0	1	0	0
0	0	1	0	0	1
0	0	0	0	1	0

convolution  
stride=1

1	1	1	1
0	1	1	2
1	1	3	0
0	2	0	0


pooling  
stride=2

1	2
2	3

# CONVNETS



## CNN architecture





# FLOWER CLASSIFICATION

# CONVNETS

## Flower Classification

daisy



# CONVNETS

## Flower Classification

dandelion





# CONVNETS

## Flower Classification

rose



# CONVNETS

## Flower Classification

sunflower





# CONVNETS

## Flower Classification

tulip



# CONVNETS



[http://download.tensorflow.org/example\\_images/flower\\_photos.tgz](http://download.tensorflow.org/example_images/flower_photos.tgz) (3,670)

## Tiny version

[https://www.visioncog.com/rpk/tiny\\_FR.zip](https://www.visioncog.com/rpk/tiny_FR.zip) (500)

### Flower Classification

(100 each, size of image varies)

- daisy
- dandelion
- rose
- sunflower
- tulip



# CONVNETS



*# Original dataset*

*# [http://download.tensorflow.org/example\\_images/flower\\_photos.tgz](http://download.tensorflow.org/example_images/flower_photos.tgz)*

*# Download tiny version of the dataset from VisionCog website*

*# After download and unzip, remember to comment the following two lines.*


```
!wget https://www.visioncog.com/rpk/tiny_FR.zip
```




```
!unzip tiny_FR.zip
```

# CONVNETS



Table of contents   Code snippets   **Files** ✕

⬆️ UPLOAD   ↻ REFRESH    MOUNT DRIVE

- ⬆️ ..
- ▶️  sample\_data
- ▶️  tiny\_FR
-  tiny\_FR.zip

# CONVNETS



*# Original dataset*

*# [http://download.tensorflow.org/example\\_images/flower\\_photos.tgz](http://download.tensorflow.org/example_images/flower_photos.tgz)*

*# Download tiny version of the dataset from VisionCog website*

*# After download and unzip, remember to comment the following two lines.*

*#!/wget [https://www.visioncog.com/rpk/tiny\\_FR.zip](https://www.visioncog.com/rpk/tiny_FR.zip)*

*#!/unzip tiny\_FR.zip*

VisionCog Research and Development Pvt. Ltd.

# CONVNETS



```
# Install TensorFlow
try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)

# TensorFlow 2.x selected.
# 2.0.0-rc1
```



# CONVNETS



```
from tensorflow import keras  
tf.random.set_seed(42)
```

```
import numpy as np  
np.random.seed(42)
```

```
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
import glob  
import PIL  
from PIL import Image
```

# CONVNETS



```
imgFiles = glob.glob("tiny_FR/*/*.jpg")
```

```
for items in imgFiles[:5]:  
    print(items)
```

```
# tiny_FR/sunflower/1715303025_e7065327e2.jpg  
# tiny_FR/sunflower/2442985637_8748180f69.jpg  
# tiny_FR/sunflower/27466794_57e4fe5656.jpg  
# tiny_FR/sunflower/40411019_526f3fc8d9_m.jpg  
# tiny_FR/sunflower/253586685_ee5b5f5232.jpg
```



# CONVNETS



```
X = []
y = []

for fName in imgFiles:

    X_i = Image.open(fName) # tiny_FR/sunflower/1715303025_e7065327e2.jpg (500, 333)

    X_i = X_i.resize((299,299)) # To make them appropriate to Xception model when using Transfer Learning

    X_i = np.array(X_i) / 255.0 # Normalize to range 0.0 to 1.0 (not stretching, only scaling)

    X.append(X_i)

    label = fName.split("/") # ['tiny_FR', 'sunflower', '1715303025_e7065327e2.jpg']

    y_i = label[1] # 'sunflower'

    y.append(y_i)
```

# CONVNETS



```
print(set(y))  
# {'daisy', 'sunflower', 'dandelion', 'rose', 'tulip'}
```

```
from sklearn.preprocessing import LabelEncoder
```

```
lEncoder = LabelEncoder()  
y = lEncoder.fit_transform(y)
```

```
print(set(y))  
# {0, 1, 2, 3, 4}
```

```
print(lEncoder.classes_)  
# ['daisy' 'dandelion' 'rose' 'sunflower' 'tulip']
```

```
y = np.array(y)
```

```
print(y.shape)
# (500,)
```

```
from sklearn.model_selection import train_test_split
```

[illegible]

# CONVNETS



```
print("X_train_shape: {}".format(X_train.shape))
```

```
# X_train_shape: (400, 299, 299, 3)
```

```
print("X_test_shape: {}".format(X_test.shape))
```

```
# X_test_shape: (100, 299, 299, 3)
```

```
# Standard scaling
```

```
mu = X_train.mean()
```

```
std = X_train.std()
```

```
X_train_std = (X_train-mu)/std
```

```
X_test_std = (X_test-mu)/std
```

# CONVNETS



*# Create the network using Functional API method*

```
input_ = keras.layers.Input(shape = X_train.shape[1:])

x = keras.layers.Conv2D(filters=32, kernel_size=5, padding='same', activation='relu')(input_)
x = keras.layers.MaxPool2D(pool_size=2)(x)
x = keras.layers.Conv2D(filters=64, kernel_size=3, padding='same', activation='relu')(x)
x = keras.layers.MaxPool2D(pool_size=2)(x)

x = keras.layers.Flatten()(x)
x = keras.layers.Dense(units=1000, activation='relu')(x)
x = keras.layers.Dense(units=100, activation='relu')(x)

output_ = keras.layers.Dense(units=5, activation='softmax')(x)

model_CNN = keras.models.Model(inputs=[input_], outputs=[output_])
```

# CONVNETS



```
model_CNN.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 299, 299, 3)]	0
conv2d (Conv2D)	(None, 299, 299, 32)	2432
max_pooling2d (MaxPooling2D)	(None, 149, 149, 32)	0
conv2d_1 (Conv2D)	(None, 149, 149, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 64)	0
flatten (Flatten)	(None, 350464)	0
dense (Dense)	(None, 1000)	350465000
dense_1 (Dense)	(None, 100)	100100
dense_2 (Dense)	(None, 5)	505
Total params: 350,586,533		
Trainable params: 350,586,533		
Non-trainable params: 0		



# CONVNETS



```
model_CNN.compile(loss='sparse_categorical_crossentropy',  
                  optimizer='adam', metrics=['accuracy'])
```

```
history_CNN = model_CNN.fit(x = X_train_std, y = y_train, epochs=25,  
                           validation_split=0.1, batch_size=32)
```

Train on 360 samples, validate on 40 samples

Epoch 1/25

360/360 [===]-7s 18ms/sample - loss: 27.5400 - accuracy: 0.2361 - val\_loss: 2.1541 - val\_accuracy: 0.0750

Epoch 2/25

360/360 [===]-2s 6ms/sample - loss: 1.4595 - accuracy: 0.3083 - val\_loss: 1.5361 - val\_accuracy: 0.2250

Epoch 3/25

360/360 [===]-2s 6ms/sample - loss: 1.3308 - accuracy: 0.5194 - val\_loss: 1.4175 - val\_accuracy: 0.4250

...

Epoch 24/25

360/360 [===]-2s 6ms/sample - loss: 1.1731e-04 - accuracy: 1.0000 - val\_loss: 3.4437 - val\_accuracy: 0.3500

Epoch 25/25

360/360 [===]-2s 6ms/sample - loss: 1.0730e-04 - accuracy: 1.0000 - val\_loss: 3.4645 - val\_accuracy: 0.3500

# CONVNETS



Train on 360 samples, validate on 40 samples

Epoch 1/25

360/360 [===]-7s 18ms/sample - loss: 27.5400 - accuracy: 0.2361 - val\_loss: 2.1541 - val\_accuracy: 0.0750

Epoch 2/25

360/360 [===]-2s 6ms/sample - loss: 1.4595 - accuracy: 0.3083 - val\_loss: 1.5361 - val\_accuracy: 0.2250

Epoch 3/25

360/360 [===]-2s 6ms/sample - loss: 1.3308 - accuracy: 0.5194 - val\_loss: 1.4175 - val\_accuracy: 0.4250

...

Epoch 24/25

360/360 [===]-2s 6ms/sample - loss: 1.1731e-04 - accuracy: 1.0000 - val\_loss: 3.4437 - val\_accuracy: 0.3500

Epoch 25/25

360/360 [===]-2s 6ms/sample - loss: 1.0730e-04 - accuracy: 1.0000 - val\_loss: 3.4645 - val\_accuracy: 0.3500



# CONVNETS




```
keys = ['accuracy', 'val_accuracy']  
progress = {k:v for k,v in history_CNN.history.items() if k in keys}
```

```
import pandas as pd  
pd.DataFrame(progress).plot()
```

```
plt.xlabel("epochs")  
plt.ylabel("accuracy")
```

```
plt.grid(True)  
plt.show()
```



# CONVNETS



```
test_loss, test_accuracy = model_CNN.evaluate(X_test_std, y_test)
# 100/1 [===] - 0s 3ms/sample - loss: 5.2441 - accuracy: 0.5500

print("Test-loss: %f, Test-accuracy: %f" % (test_loss, test_accuracy))
# Test-loss: 2.738503, Test-accuracy: 0.550000
```

VisionCog Research and Development Pvt. Ltd.