



TRANSFER LEARNING

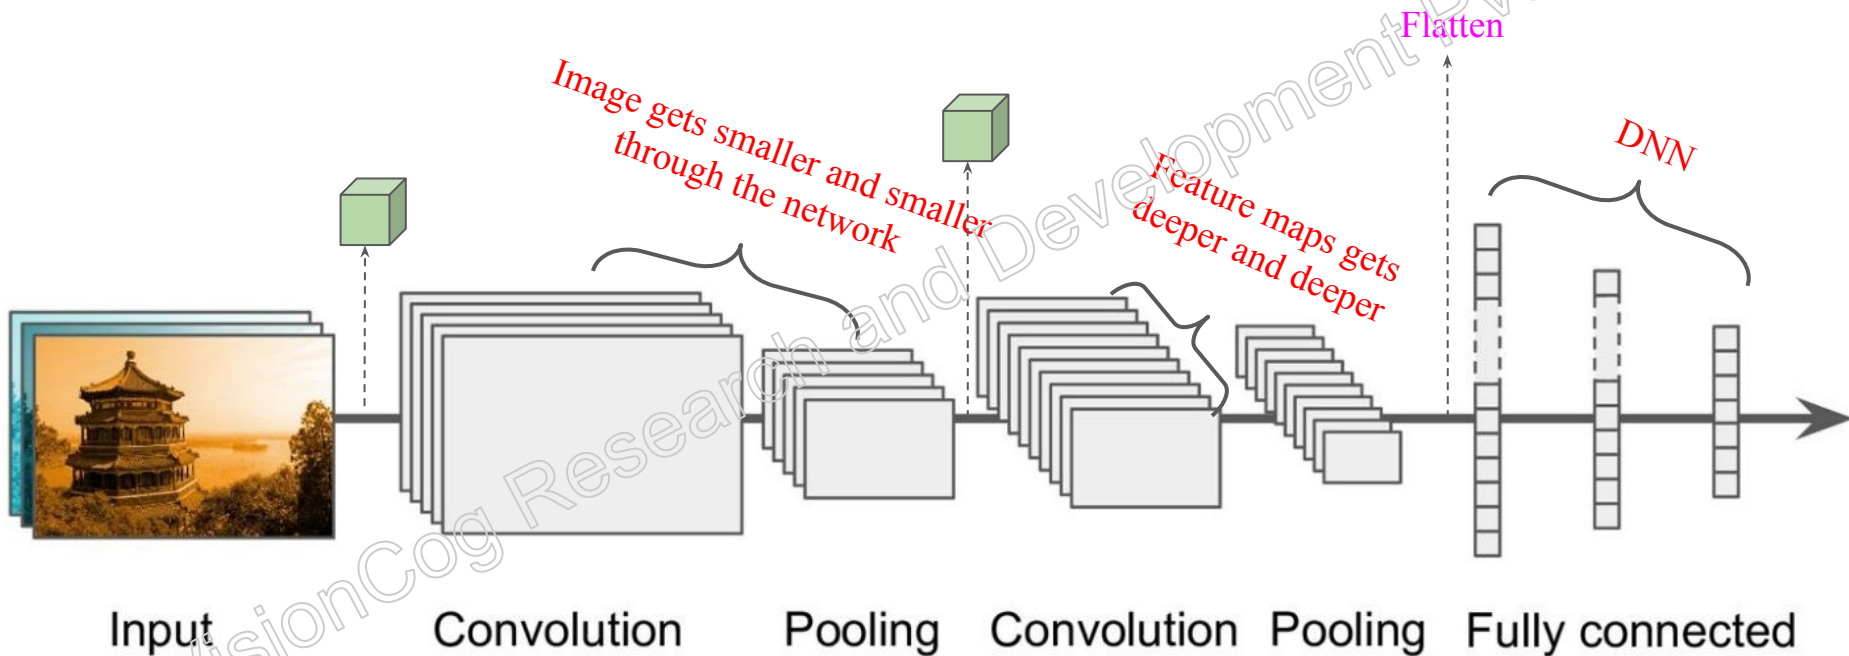
Dr. Ram Prasad K
VisionCog R&D

ram.krish@visioncog.com
<https://www.visioncog.com>

CONVNETS




CNN architecture



TRANSFER LEARNING

CONVNETS


Transfer learning



CONVNETS



Transfer learning and Fine-Tuning





Xception: Deep Learning with Depthwise Separable Convolutions

François Chollet

Google, Inc.

fchollet@google.com

Abstract

We present an interpretation of Inception modules in convolutional neural networks as being an intermediate step in-between regular convolution and the depthwise separable convolution operation (a depthwise convolution followed by a pointwise convolution). In this light, a depthwise separable convolution can be understood as an Inception module with a maximally large number of towers. This observation leads us to propose a novel deep convolutional neural network architecture inspired by Inception, where Inception modules have been replaced with depthwise separable convolutions. We show that this architecture, dubbed Xception, slightly outperforms Inception V3 on the ImageNet dataset (which Inception V3 was designed for), and significantly outperforms Inception V3 on a larger image classification dataset comprising 350 million images and 17,000 classes. Since the Xception architecture has the same number of parameters as Inception V3, the performance gains are not due to increased capacity but rather to a more efficient use of model parameters.

as GoogLeNet (Inception V1), later refined as Inception V2 [7], Inception V3 [21], and most recently Inception-ResNet [19]. Inception itself was inspired by the earlier Network-In-Network architecture [11]. Since its first introduction, Inception has been one of the best performing family of models on the ImageNet dataset [14], as well as internal datasets in use at Google, in particular JFT [5].

The fundamental building block of Inception-style models is the Inception module, of which several different versions exist. In figure 1 we show the canonical form of an Inception module, as found in the Inception V3 architecture. An Inception model can be understood as a stack of such modules. This is a departure from earlier VGG-style networks which were stacks of simple convolution layers.

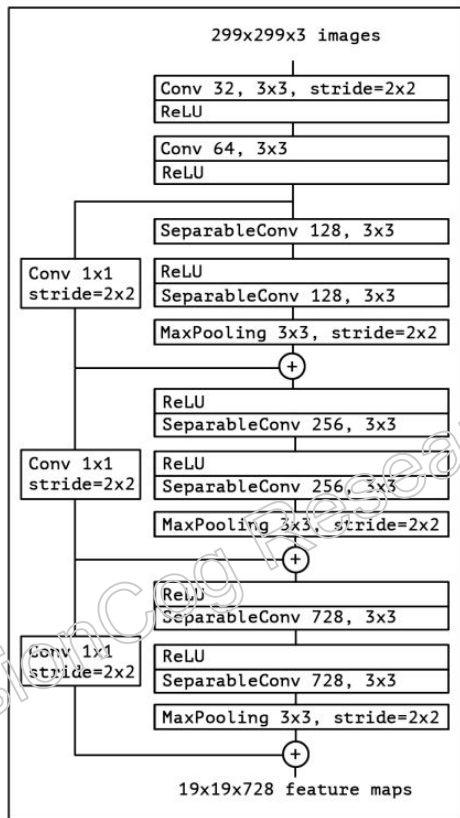
While Inception modules are conceptually similar to convolutions (they are convolutional feature extractors), they empirically appear to be capable of learning richer representations with less parameters. How do they work, and how do they differ from regular convolutions? What design strategies come after Inception?

1.1. The Inception hypothesis

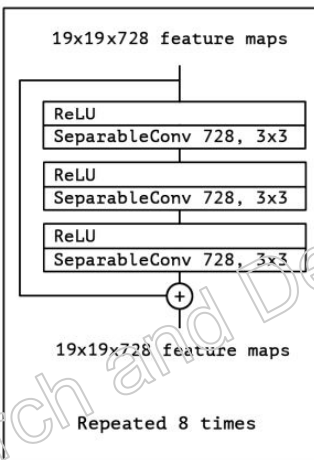
CONVNETS



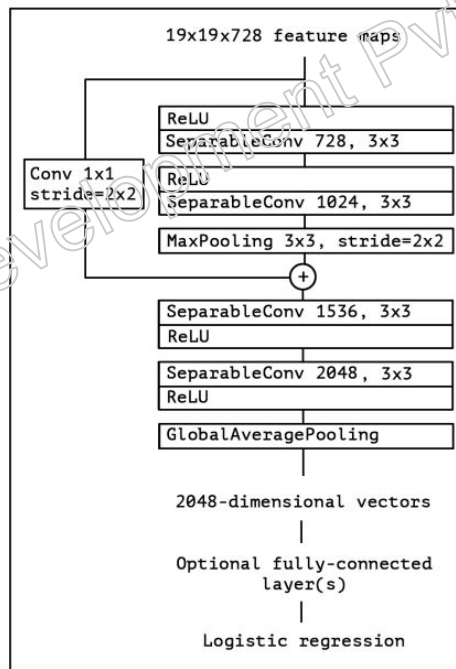
Entry flow



Middle flow



Exit flow





FLOWER CLASSIFICATION

CONVNETS

Flower Classification

daisy



CONVNETS

Flower Classification

dandelion



CONVNETS

Flower Classification

rose



VisionCog Research and Development Pvt. Ltd.

CONVNETS

Flower Classification

sunflower



CONVNETS

Flower Classification

tulip



CONVNETS



http://download.tensorflow.org/example_images/flower_photos.tgz (3,670)

Tiny version

https://www.visioncog.com/rpk/tiny_FR.zip (500)

Flower Classification

(100 each, size of image varies)

- daisy
- dandelion
- rose
- sunflower
- tulip



CONVNETS



Original dataset

http://download.tensorflow.org/example_images/flower_photos.tgz

Download tiny version of the dataset from VisionCog website

After download and unzip, remember to comment the following two lines.


```
!wget https://www.visioncog.com/rpk/tiny_FR.zip
```

```
!unzip tiny_FR.zip
```

CONVNETS



Table of contents Code snippets **Files** ✕

⬆️ UPLOAD ↻ REFRESH  MOUNT DRIVE

- ⬆️ ..
- ▶️ sample_data
- ▶️ tiny_FR
- 📄 tiny_FR.zip

CONVNETS



Original dataset

http://download.tensorflow.org/example_images/flower_photos.tgz

Download tiny version of the dataset from VisionCog website

After download and unzip, remember to comment the following two lines.

#!/wget https://www.visioncog.com/rpk/tiny_FR.zip

#!/unzip tiny_FR.zip

VisionCog Research and Development Pvt. Ltd.

CONVNETS



```
# Install TensorFlow
try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass

import tensorflow as tf
print(tf.__version__)

# TensorFlow 2.x selected.
# 2.0.0-rc1
```

CONVNETS



```
from tensorflow import keras  
tf.random.set_seed(42)
```

```
import numpy as np  
np.random.seed(42)
```

```
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
import glob  
import PIL  
from PIL import Image
```

CONVNETS



```
imgFiles = glob.glob("tiny_FR/*/*.jpg")
```

```
for items in imgFiles[:5]:  
    print(items)
```

```
# tiny_FR/sunflower/1715303025_e7065327e2.jpg  
# tiny_FR/sunflower/2442985637_8748180f69.jpg  
# tiny_FR/sunflower/27466794_57e4fe5656.jpg  
# tiny_FR/sunflower/40411019_526f3fc8d9_m.jpg  
# tiny_FR/sunflower/253586685_ee5b5f5232.jpg
```

CONVNETS



```
X = []
y = []

for fName in imgFiles:

    X_i = Image.open(fName) # tiny_FR/sunflower/1715303025_e7065327e2.jpg (500, 333)

    X_i = X_i.resize((299,299)) # To make them appropriate to Xception model when using Transfer Learning

    X_i = np.array(X_i) / 255.0 # Normalize to range 0.0 to 1.0 (not stretching, only scaling)

    X.append(X_i)

    label = fName.split("/") # ['tiny_FR', 'sunflower', '1715303025_e7065327e2.jpg']

    y_i = label[1] # 'sunflower'

    y.append(y_i)
```

CONVNETS



```
print(set(y))  
# {'daisy', 'sunflower', 'dandelion', 'rose', 'tulip'}
```

```
from sklearn.preprocessing import LabelEncoder
```

```
lEncoder = LabelEncoder()  
y = lEncoder.fit_transform(y)
```

```
print(set(y))  
# {0, 1, 2, 3, 4}
```

```
print(lEncoder.classes_)  
# ['daisy' 'dandelion' 'rose' 'sunflower' 'tulip']
```

```
y = np.array(y)
```

```
print(y.shape)
# (500,)
```

```
from sklearn.model_selection import train_test_split
```

[illegible]

CONVNETS



```
print("X_train_shape: {}".format(X_train.shape))
```

```
# X_train_shape: (400, 299, 299, 3)
```

```
print("X_test_shape: {}".format(X_test.shape))
```

```
# X_test_shape: (100, 299, 299, 3)
```

```
# Standard scaling
```

```
mu = X_train.mean()
```

```
std = X_train.std()
```

```
X_train_std = (X_train-mu)/std
```

```
X_test_std = (X_test-mu)/std
```


CONVNETS

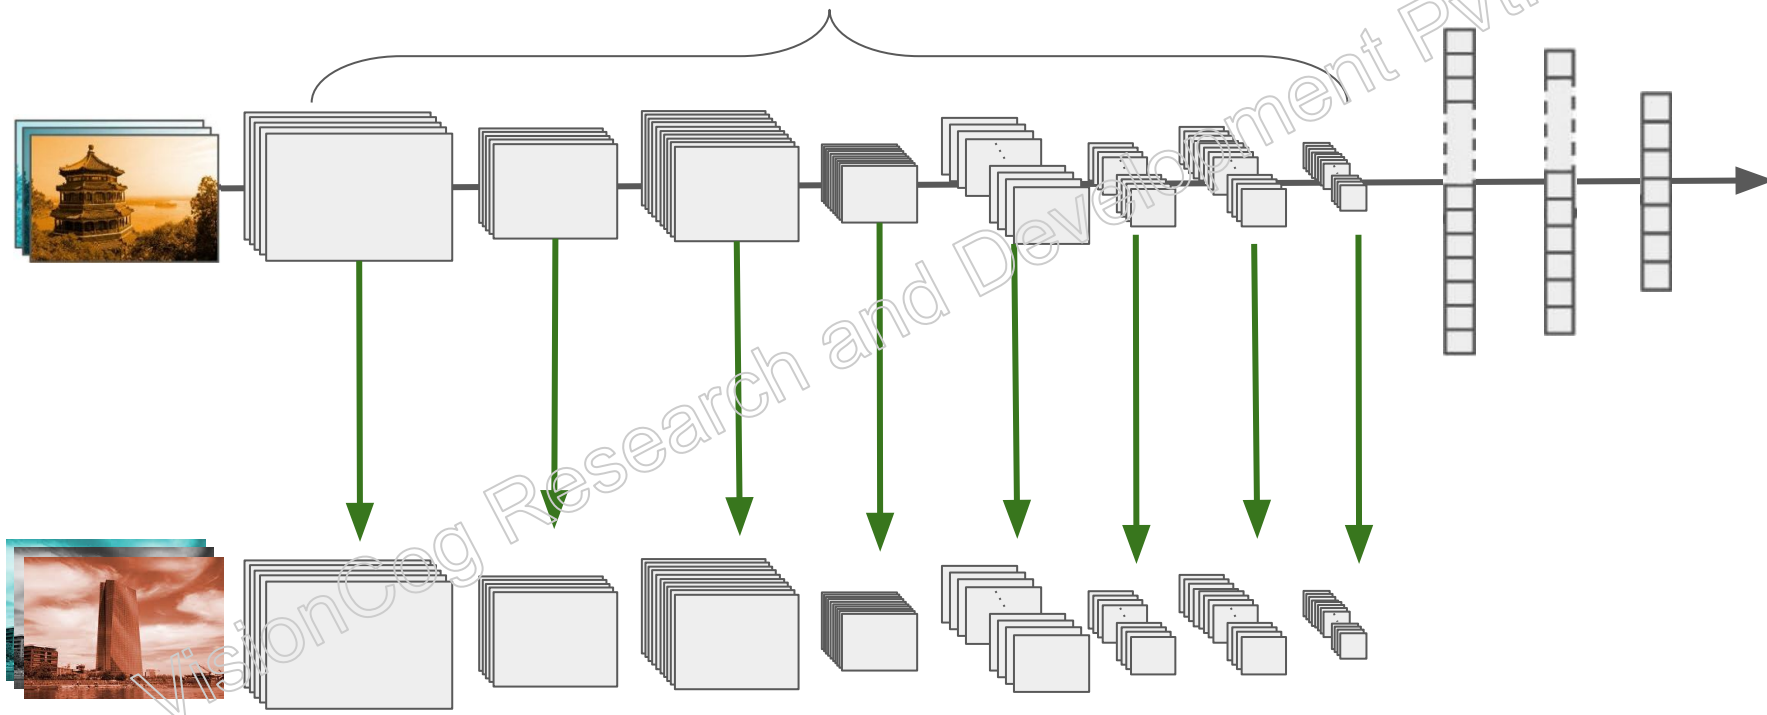


```
base_model = keras.applications.xception.Xception(weights='imagenet',  
                                                    include_top=False)
```

VisionCog Research and Development Pvt. Ltd.

CONVNETS

Transfer learning



CONVNETS



```
base_model = keras.applications.xception.Xception(weights='imagenet',  
                                                    include_top=False)
```


```
for layer in base_model.layers:  
    layer.trainable = False
```

```
global_pool = keras.layers.GlobalAveragePooling2D()(base_model.output)  
output_ = keras.layers.Dense(units=5, activation='softmax')(global_pool)
```

```
model_TL = keras.models.Model(inputs=[base_model.input], outputs=[output_])
```

CONVNETS

Transfer learning



CONVNETS



```
model_TL.compile(loss='sparse_categorical_crossentropy',  
                 optimize='adam', metrics=['accuracy'])
```

```
history_TL = model_TL.fit(x = X_train_std, y = y_train, epochs=25,  
                          validation_split=0.1, batch_size=16)
```

Train on 360 samples, validate on 40 samples

Epoch 1/25

360/360 [===]-27s 75ms/sample - loss: 0.9786 - accuracy: 0.6500 - val_loss: 7.0164 - val_accuracy: 0.5500

Epoch 2/25

360/360 [===]-15s 40ms/sample - loss: 0.4938 - accuracy: 0.8583 - val_loss: 10.1537 - val_accuracy: 0.3000

Epoch 3/25

360/360 [===]-14s 40ms/sample - loss: 0.3327 - accuracy: 0.8889 - val_loss: 6.9043 - val_accuracy: 0.4000

...

Epoch 24/25

360/360 [===]-14s 39ms/sample - loss: 0.0482 - accuracy: 0.9806 - val_loss: 1.0079 - val_accuracy: 0.7500

Epoch 25/25

360/360 [===]-14s 39ms/sample - loss: 0.0650 - accuracy: 0.9750 - val_loss: 7.1623 - val_accuracy: 0.6250

CONVNETS



Train on 360 samples, validate on 40 samples

Epoch 1/25

360/360 [===]-27s 75ms/sample - loss: 0.9786 - accuracy: 0.6500 - val_loss: 7.0164 - val_accuracy: 0.5500

Epoch 2/25

360/360 [===]-15s 40ms/sample - loss: 0.4938 - accuracy: 0.8583 - val_loss: 10.1537 - val_accuracy: 0.3000

Epoch 3/25

360/360 [===]-14s 40ms/sample - loss: 0.3327 - accuracy: 0.8889 - val_loss: 6.9043 - val_accuracy: 0.4000

...

Epoch 24/25

360/360 [===]-14s 39ms/sample - loss: 0.0482 - accuracy: 0.9806 - val_loss: 1.0079 - val_accuracy: 0.7500

Epoch 25/25

360/360 [===]-14s 39ms/sample - loss: 0.0650 - accuracy: 0.9750 - val_loss: 7.1623 - val_accuracy: 0.6250

CONVNETS



```
testLoss_TL, testAccuracy_TL = model_TL.evaluate(x = X_test_std, y = y_test)
# 100/1 [===] - 3s 26ms/sample - loss: 2.6403 - accuracy: 0.7200

print("Test-loss: %f, Test-accuracy: %f" % (testLoss_TL, testAccuracy_TL))
# Test-loss: 3.964784, Test-accuracy: 0.720000
```